

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
import seaborn as sns
from sklearn.linear_model import Ridge, RidgeCV, Lasso
from sklearn.preprocessing import StandardScaler
```

```
In [2]: df=pd.read_csv(r"C:\Users\Y.Saranya\Downloads\fiat500_VehicleSelection_Dataset")
df
```

```
Out[2]:
```

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon
0	1	lounge	51	882	25000	1	44.907242	8.611560
1	2	pop	51	1186	32500	1	45.666359	12.241890
2	3	sport	74	4658	142228	1	45.503300	11.417840
3	4	lounge	51	2739	160000	1	40.633171	17.634609
4	5	pop	73	3074	106880	1	41.903221	12.495650
...
1533	1534	sport	51	3712	115280	1	45.069679	7.704920
1534	1535	lounge	74	3835	112000	1	45.845692	8.666870
1535	1536	pop	51	2223	60457	1	45.481541	9.413480
1536	1537	lounge	51	2557	80750	1	45.000702	7.682270
1537	1538	pop	51	1766	54276	1	40.323410	17.568270

1538 rows × 9 columns



In [3]: `df.head(10)`

Out[3]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	lounge	51	882	25000	1	44.907242	8.611560	8900
1	2	pop	51	1186	32500	1	45.666359	12.241890	8800
2	3	sport	74	4658	142228	1	45.503300	11.417840	4200
3	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
4	5	pop	73	3074	106880	1	41.903221	12.495650	5700
5	6	pop	74	3623	70225	1	45.000702	7.682270	7900
6	7	lounge	51	731	11600	1	44.907242	8.611560	10700
7	8	lounge	51	1521	49076	1	41.903221	12.495650	9100
8	9	sport	73	4049	76000	1	45.548000	11.549470	5600
9	10	sport	51	3653	89000	1	45.438301	10.991700	6000

In [4]: `df.info`

Out[4]:

```
<bound method DataFrame.info of
ID      model  engine_power  age_in_days
0         1   lounge         51          882    25000
1         2     pop         51         1186    32500
2         3   sport         74         4658   142228
3         4   lounge         51         2739   160000
4         5     pop         73         3074   106880
...      ...      ...      ...      ...
1533  1534   sport         51         3712   115280
1534  1535   lounge         74         3835   112000
1535  1536     pop         51         2223    60457
1536  1537   lounge         51         2557    80750
1537  1538     pop         51         1766    54276

lat      lon  price
0    44.907242  8.611560  8900
1    45.666359 12.241890  8800
2    45.503300 11.417840  4200
3    40.633171 17.634609  6000
4    41.903221 12.495650  5700
...      ...      ...      ...
1533  45.069679  7.704920  5200
1534  45.845692  8.666870  4600
1535  45.481541  9.413480  7500
1536  45.000702  7.682270  5990
1537  40.323410 17.568270  7900

[1538 rows x 9 columns]>
```

In [5]: `df.describe`

```
Out[5]: <bound method NDFrame.describe of
ys      km  previous_owners \
0        1    lounge        51      882    25000      1
1        2      pop        51     1186    32500      1
2        3    sport        74     4658   142228      1
3        4    lounge        51     2739   160000      1
4        5      pop        73     3074   106880      1
...     ...     ...     ...     ...     ...     ...
1533  1534    sport        51     3712   115280      1
1534  1535    lounge        74     3835   112000      1
1535  1536      pop        51     2223    60457      1
1536  1537    lounge        51     2557    80750      1
1537  1538      pop        51     1766    54276      1

      lat      lon  price
0  44.907242  8.611560  8900
1  45.666359  12.241890  8800
2  45.503300  11.417840  4200
3  40.633171  17.634609  6000
4  41.903221  12.495650  5700
...     ...     ...     ...
1533  45.069679  7.704920  5200
1534  45.845692  8.666870  4600
1535  45.481541  9.413480  7500
1536  45.000702  7.682270  5990
1537  40.323410  17.568270  7900


[1538 rows x 9 columns]>
```

In [6]: `df.columns`

```
Out[6]: Index(['ID', 'model', 'engine_power', 'age_in_days', 'km', 'previous_owners',
              'lat', 'lon', 'price'],
              dtype='object')
```


```
In [7]: sns.displot(['Price'])
```

```
Out[7]: <seaborn.axisgrid.FacetGrid at 0x23fc6794490>
```




```
In [8]: sns.displot(df['price'])
```

```
Out[8]: <seaborn.axisgrid.FacetGrid at 0x23fcca264d0>
```




```
In [9]: sns.displot(df['engine_power'])
```

```
Out[9]: <seaborn.axisgrid.FacetGrid at 0x23fcd165f30>
```




```
In [10]: fiatdf=df[['ID', 'engine_power', 'age_in_days', 'km', 'previous_owners',  
                  'lat', 'lon','price']]  
sns.heatmap(fiatdf.corr())#with price
```

Out[10]: <Axes: >



```
In [11]: fiatdf=df[['ID', 'engine_power', 'age_in_days', 'km', 'previous_owners',  
                'lat', 'lon']]  
sns.heatmap(fiatdf.corr())#without price
```

Out[11]: <Axes: >



```
In [12]: X=fiatdf[['ID', 'engine_power', 'age_in_days', 'km', 'previous_owners',  
                'lat', 'lon']]  
y=df['price']
```

```
In [13]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=  
regr=LinearRegression()  
regr.fit(X_train,y_train)  
print(regr.intercept_)
```

8971.195683499936



```
In [14]: coeff_df=pd.DataFrame(regr.coef_,X.columns,columns=['coefficient'])  
coeff_df
```

```
Out[14]:
```

	coefficient
ID	-0.046704
engine_power	11.646408
age_in_days	-0.898018
km	-0.017232
previous_owners	26.400886
lat	32.189709
lon	0.161073


```
In [15]: predictions=regr.predict(X_test)  
plt.scatter(y_test,predictions)
```

```
Out[15]: <matplotlib.collections.PathCollection at 0x23fce477ac0>
```



```
In [16]: sns.displot((y_test,predictions),bins=50)
```

```
Out[16]: <seaborn.axisgrid.FacetGrid at 0x23fce36ace0>
```



```
In [17]: from sklearn import metrics
print('MAE:',metrics.mean_absolute_error(y_test,predictions))
print('MSE:',metrics.mean_squared_error(y_test,predictions))
print('MAE:',np.sqrt(metrics.mean_squared_error(y_test,predictions)))
```

MAE: 593.0876179519935

MSE: 551442.6799691805


MAE: 742.5918663500029

```
In [18]: #accuracy
regr=LinearRegression()
regr.fit(X_train,y_train)
regr.fit(X_train,y_train)
print(regr.score(X_test,y_test))
```

0.8597136704308866


```
In [19]: plt.figure(figsize=(10,10))  
sns.heatmap(fiatdf.corr(),annot=True)
```

Out[19]: <Axes: >




```
In [20]: sns.lmplot(x="age_in_days",y="km",data=fiatdf,order=2)
```

```
Out[20]: <seaborn.axisgrid.FacetGrid at 0x23fceaa20b0>
```



```
In [21]: sns.lmplot(x="price",y="km",data=df,order=2)
```

```
Out[21]: <seaborn.axisgrid.FacetGrid at 0x23fceb02bf0>
```



```
In [22]: df.fillna(method='ffill',inplace=True)
x=np.array(df['age_in_days']).reshape(-1,1)
y=np.array(df['km']).reshape(-1,1)
df.dropna(inplace=True)
```


```
In [23]: X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
regr.fit(X_train,y_train)
regr.fit(X_train,y_train)
```

```
Out[23]: LinearRegression()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [24]: y_pred=regr.predict(X_test)
plt.scatter(X_test,y_test,color='y')
plt.plot(X_test,y_pred,color='b')
plt.show()
```



```
In [25]: #Linear regression model
regr=LinearRegression()
regr.fit(X_train,y_train)
actual=y_test #actual value
train_score_regr=regr.score(X_train,y_train)
test_score_regr=regr.score(X_test,y_test)
print("\nLinear model:\n")
print("The train score for Linear model is {}".format(train_score_regr))
print("The test score for Linear model is {}".format(test_score_regr))
```

Linear model:

The train score for Linear model is 0.698875263814575

The test score for Linear model is 0.6788811761009449

```
In [26]: #ridge regression model
ridgeReg=Ridge(alpha=10)
ridgeReg.fit(X_train,y_train)
#train and test score for ridge regression
train_score_ridge=ridgeReg.score(X_train,y_train)
test_score_ridge=ridgeReg.score(X_test,y_test)
print("\nRidge model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge model:

The train score for ridge model is 0.698875263814575

The test score for ridge model is 0.6788811770230736

```
In [27]: #Lasso regression model
lassoReg=Lasso(alpha=10)
lassoReg.fit(X_train,y_train)
#train and test score for ridge regression
train_score_lasso=lassoReg.score(X_train,y_train)
test_score_lasso=lassoReg.score(X_test,y_test)
print("\nLasso model:\n")
print("The train score for lasso model is {}".format(train_score_lasso))
print("The test score for lasso model is {}".format(test_score_lasso))
```

Lasso model:

The train score for lasso model is 0.6988752638145399

The test score for lasso model is 0.6788812133066009

```
In [28]: #using the linear cv model for ridge regression
from sklearn.linear_model import RidgeCV
#ridge cross validation
ridge_cv=RidgeCV(alphas=[0.0001,0.001,0.01,0.1,1,10]).fit(X_train,y_train)
#score
print(ridge_cv.score(X_train,y_train))
print(ridge_cv.score(X_test,y_test))
```

0.6988752638145734

0.6788811836133852

```
In [30]: #using the linear cv model for lasso regression  
from sklearn.linear_model import LassoCV  
#Lasso cross validation  
lasso_cv=LassoCV(alphas=[0.0001,0.001,0.01,0.1,1,10],random_state=0).fit(X_train,y_train)  
#score  
print(lasso_cv.score(X_train,y_train))  
print(lasso_cv.score(X_test,y_test))
```

0.698875263814575

0.6788811761013169

C:\Users\Y.Saranya\anaconda3\lib\site-packages\sklearn\linear_model_coordinate_descent.py:1568: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

y = column_or_1d(y, warn=True)

In []: